

```
# coding: utf-8
```

```
### Predicting Continuous Target Variables With Regression Analysis
```

```
# - Read the dataset from the url into a data frame.
```

```
# - Display the first few rows of the data frame to make sure the data was read properly
```

```
# In[1]:
```

```
import pandas as pd
```

```
df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/housing/housing.data',  
header=None, sep='\s+')
```

```
df.columns = ['CRIM', 'ZN', 'INDUS', 'CHAS',  
              'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
              'TAX', 'PTRATIO', 'B', 'LSTAT', 'MEDV']
```

```
df.head()
```

```
# - Visualize the important characteristics of the dataset before attempting to build a regression model  
(Exploratory Data Analysis - EDA - is a recommended first step prior to the training of a machine learning  
model. We will create a scatterplot matrix that allows us to visualize the pair-wise correlations between  
different features in the dataset in one place.
```

```
# In[2]:
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
sns.set(style='whitegrid', context='notebook')
```

```
cols = ['LSTAT', 'INDUS', 'NOX', 'RM', 'MEDV']
```

```
sns.pairplot(df[cols], size=2.5)
plt.show()
```

# - Create a correlation matrix to quantify the linear relationships between features.

# In[3]:

```
import numpy as np

cor_matrix = np.corrcoef(df[cols].values.T) # note we transpose to get the data by columns. Columns
become rows.

sns.set(font_scale=1.5)

cor_heat_map = sns.heatmap(cor_matrix,
                             cbar=True,
                             annot=True,
                             square=True,
                             fmt='.2f',
                             annot_kws={'size':15},
                             yticklabels=cols,
                             xticklabels=cols)

plt.show()
```

# - Separate the independent and dependent variables into two variables X and y and also standardize the data

# In[4]:

```
X = df[['RM']].values # note how we pass a list of columns (here just a single-item list) to access data in
df
```

```
y = df[['MEDV']].values # you either have to define this list of columns separately or inline like here,
where you get the [[ ...]]
```

```
#print(df[['MEDV']].head())
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc_x = StandardScaler()
```

```
sc_y = StandardScaler()
```

```
X_std = sc_x.fit_transform(X)
```

```
y_std = sc_y.fit_transform(y)
```

```
# let's see how sklearn does the scaling
```

```
print('original data\n', y[0:4])
```

```
print('scaled data using sklearn StandardScaler\n', y_std[0:5])
```

```
#let's do the same computation manually
```

```
print('average of original data', np.average(y[:]))
```

```
print('STDV of original data', np.std(y[:]))
```

```
z_scores = (y[:] - np.average(y[:])) / np.std(y[:])
```

```
print('Manually computed Z scores:\n', z_scores[0:5] )
```

```
# - Build the linear regressionmodel
```

```
# ln[5]:
```

```
from sklearn.linear_model import LinearRegression
```

```
slr = LinearRegression()
```

```
# first, let's fit the un-standardized data
```

```
slr.fit(X,y)
```

```
print('Slope: %.3f' % slr.coef_[0])
```

```
print('Intercept: %.3f' % slr.intercept_)  
  
# now, let's try with standardized data  
  
slr_std = LinearRegression()  
slr_std.fit(X_std,y_std)  
print('Slope: %.3f' % slr_std.coef_[0])  
print('Intercept: %.3f' % slr_std.intercept_)
```

# - let's make a prediction using both models (the model using standardized data and the model using un-standardized data)

# In[6]:

# predict the price of a 5 bedroom house

```
import numpy as np
```

```
num_rooms = [5.0]
```

```
num_rooms_std = sc_x.transform(np.array(num_rooms).reshape(len(num_rooms),1)) # note transform expects a 2D array
```

```
print('standardized rooms: %.3f', num_rooms_std)
```

```
predicted_price_std = slr_std.predict(num_rooms_std)
```

```
print('Predicted Price std: %.3f' % predicted_price_std)
```

```
print("Predicted Price in $1000's using standardized data: %.3f" %  
sc_y.inverse_transform(predicted_price_std) )
```

# Now let's predict using the model that uses un-standardized data

```
predicted_price_non_std = slr.predict(np.array(num_rooms).reshape(len(num_rooms),1))
```

```
print("Predicted Price in $1000's Using un-standardized data: %.3f" % predicted_price_non_std)
```

# - Visualize how well the linear regression line fits the data

```
# - plot a scatterplot of the training data  
# - add the regression line
```

```
# In[7]:
```

```
plt.scatter(X_std,y_std, c='blue')  
plt.plot(X_std,slr_std.predict(X_std), color='red')  
plt.xlabel('Average number of rooms [RM] (standardized)')  
plt.ylabel('Price in $1000\'s [MEDV] (standardized)')  
plt.show()
```

```
# - now visualize the original, un-standardized training data and regression line
```

```
# In[8]:
```

```
plt.scatter(X,y, color='blue')  
plt.plot(X, slr.predict(X), color='red')  
plt.xlabel('Average Number of Rooms [RM]')  
plt.ylabel('Price in $1000\'s [MEDV]')  
plt.show()
```

```
# - evaluate the performance of the regression model on training data
```

```
# In[9]:
```

```
from sklearn.metrics import mean_squared_error, r2_score  
# first let's focus on original un-standardized data and model
```

```

MSE1 = mean_squared_error(y_true= y, y_pred= slr.predict(X))
r2_1 = r2_score(y_true= y, y_pred= slr.predict(X))
print('MSE for regression model using un-standrdized features: %.3f' % MSE1)
print('RMSE:', np.sqrt(MSE1))
print('r2:', r2_1)
print('=====')
#now let's do the same thing for standardized data and regression model
MSE_2 = mean_squared_error(y_true= y_std, y_pred=slr_std.predict(X_std) )
print('MSE for standardized data:', MSE_2)
y_true_converted_to_originl = sc_y.inverse_transform(y_std)
y_pred_converted_to_original = sc_y.inverse_transform(slr_std.predict(X_std))
MSE2 = mean_squared_error(y_true= y_true_converted_to_originl, y_pred=
y_pred_converted_to_original)

r2_2 = r2_score(y_true= y_std, y_pred= slr_std.predict(X_std))
r2_2_using_data_converted_to_original_scale = r2_score(y_true= y_true_converted_to_originl,
y_pred= y_pred_converted_to_original)
print('MSE using standardized features converted to original scale: %.3f' % MSE2)
print('RMSE:', np.sqrt(MSE2))
print('r2 using standardized data and model:',r2_2)
print('r2 using data converted to original scale', r2_2_using_data_converted_to_original_scale)

```